Customer Purchase Trends Pipeline

Project Overview

The Customer Purchase Trends Pipeline is a cloud-based data pipeline designed to process and analyze retail transaction data. The pipeline uses **Google Cloud Storage (GCS)** and **BigQuery** to process over 3,900 retail transactions and generate actionable insights.

The primary goal of the pipeline is to demonstrate the end-to-end flow of data from ingestion to reporting, including data cleaning, transformation, and aggregation

Architecture and Components

The pipeline follows a standard extract, transform, load (ETL) process using several Google Cloud services.

- **Ingestion (Bronze Layer)**: Daily raw CSV files are ingested and stored in a GCS bucket. This acts as the bronze layer, holding the raw, unprocessed data.
- **Transformation (Silver Layer)**: Data is transformed in BigQuery to clean missing values and standardize categories.
- **Aggregation (Gold Layer)**: The transformed data is aggregated to create fact tables in BigQuery, such as "purchases by age," "gender," and "season". These aggregated tables form the gold layer, which is optimized for analysis and reporting.
- **Visualization**: A Looker Studio dashboard is used to visualize key metrics like top product categories and seasonal sales trends.

Technical Implementation

Step 1: Project Preparation

First, several APIs must be enabled within the Google Cloud project to support the pipeline's functionality.

- **BigQuery API**: For creating, managing, and querying data.
- **BigQuery Data Transfer API**: Required for scheduled queries to automate data processing.
- **Cloud Storage JSON API**: For interacting with GCS buckets.
- **Cloud Scheduler API (Optional)**: Can be enabled later for scheduling tasks.

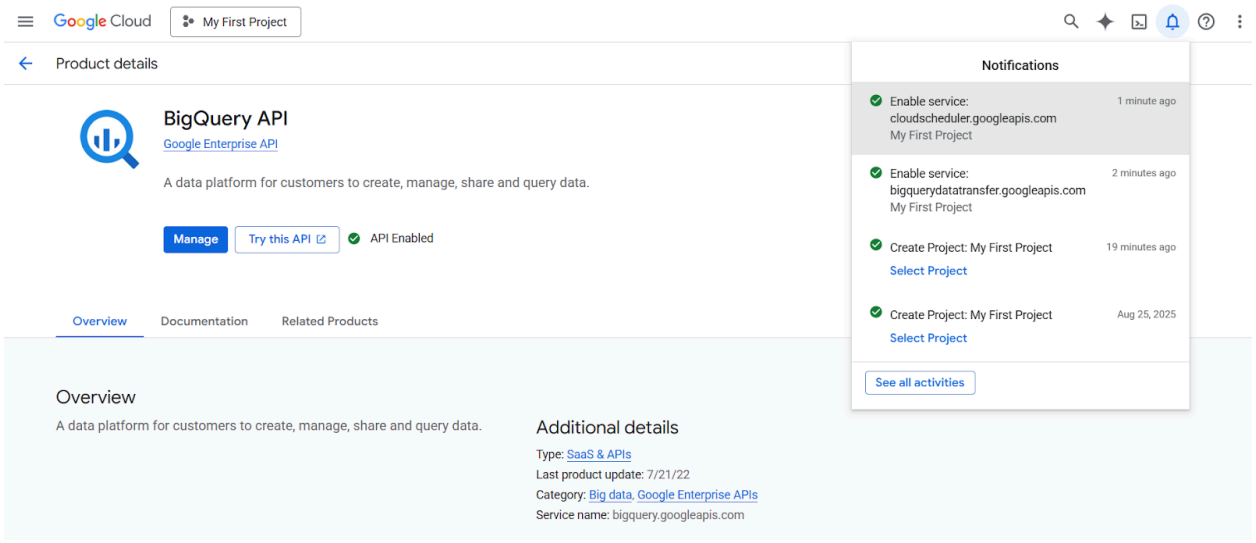

Step 2: Data Storage

A GCS bucket, named retail-pipeline-bronze, was created to store the raw daily CSV files.

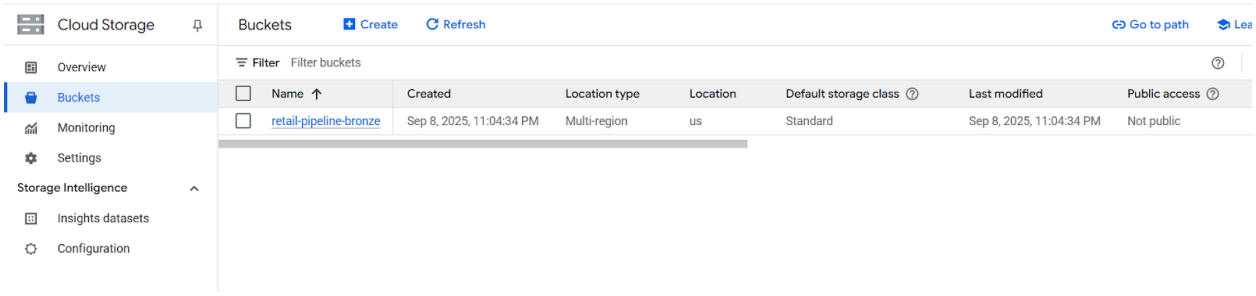

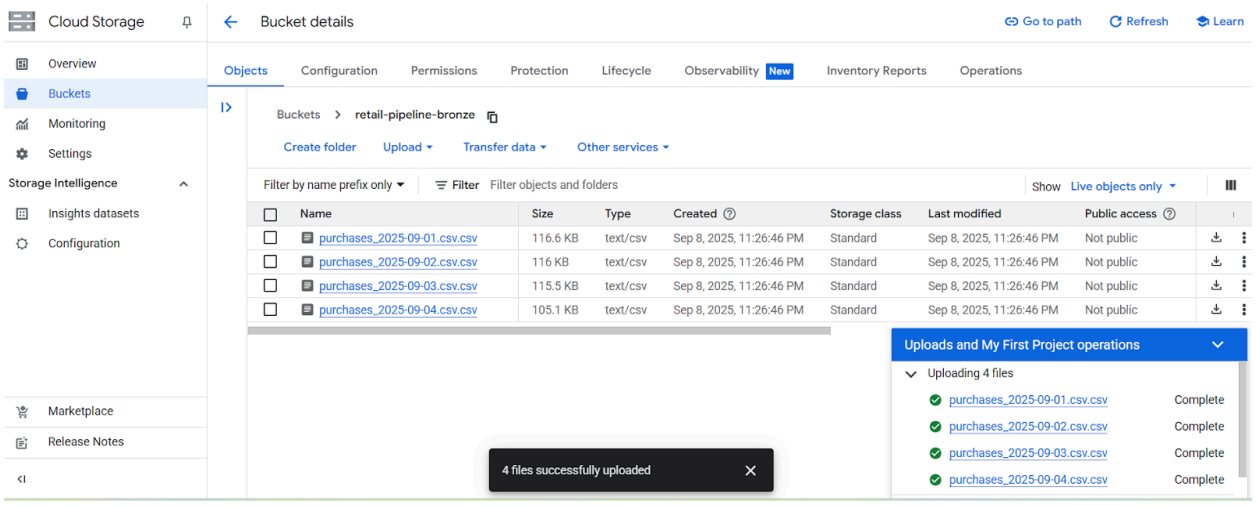

Step 3: BigQuery Setup

A BigQuery dataset, such as retail_analytics, was created to house the transformed and aggregated tables. External tables were created in BigQuery to point to the raw CSV data in GCS, allowing it to be queried directly without loading it first.

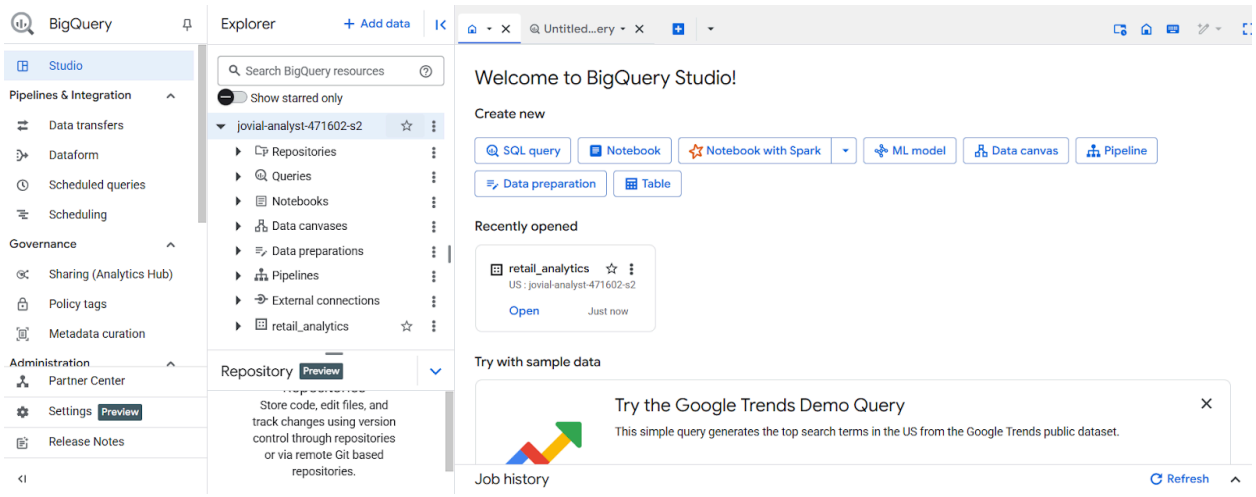

Step 4: Data Transformation and Aggregation

SQL queries were used in BigQuery to transform the raw data into aggregated fact tables. Examples include:

- Seasonal and Category Sales: A table named agg_sales_by_season_category was created by grouping purchases by season and category to calculate total sales and average order value.
- Demographic Sales: A table named agg_sales_by_age_gender was created to show sales trends segmented by age and gender.

The queries were scheduled to run on a set frequency (e.g., every 8 hours) using the BigQuery Data Transfer API to ensure the fact tables are regularly updated with new data.

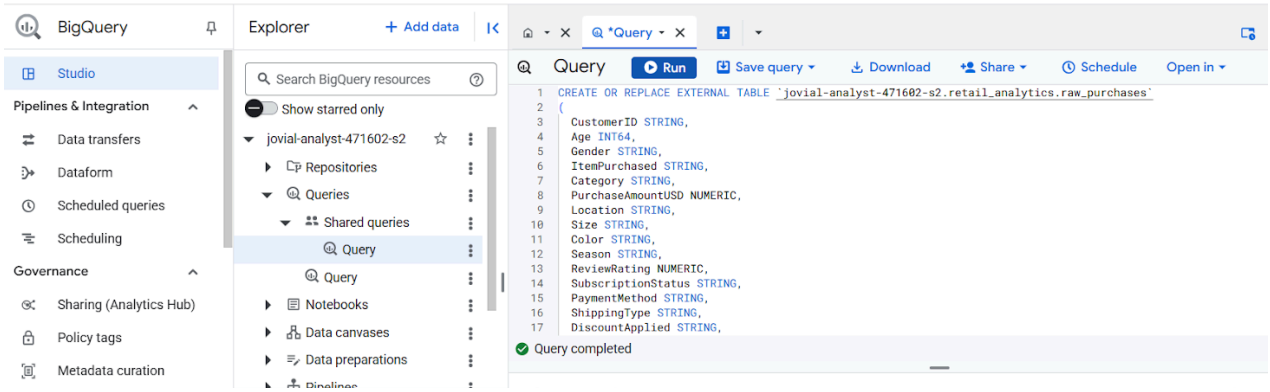

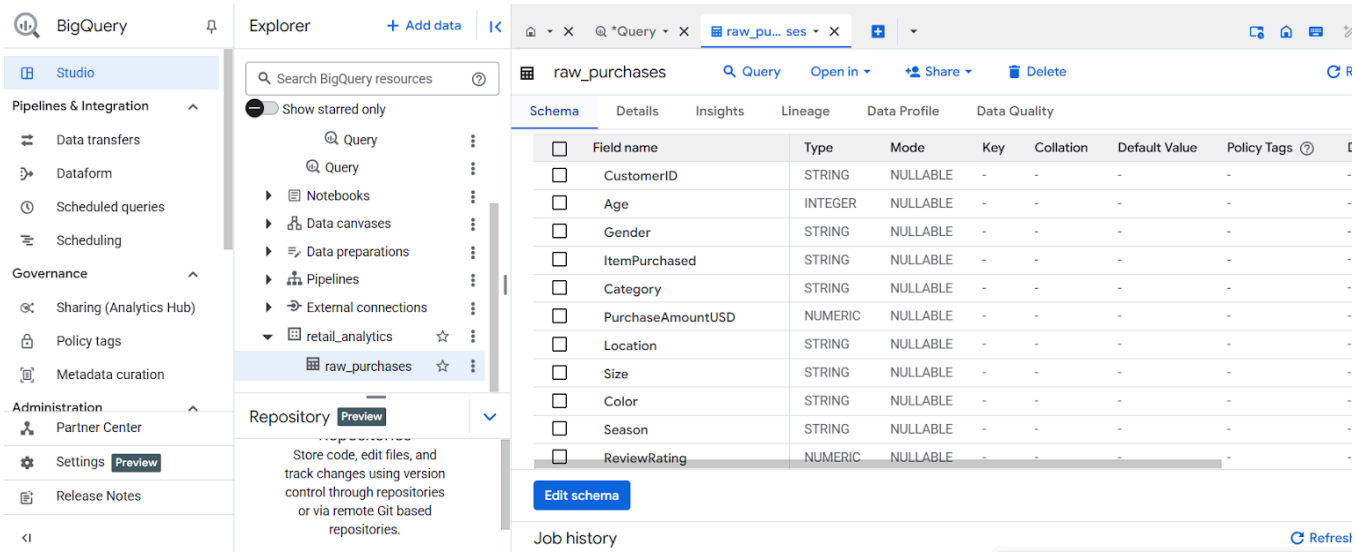
BigQuery
Explorer
+ Add data
Studio
Pipelines & Integration
Data transfers
Dataform
Scheduled queries
Scheduling
Governance
Sharing (Analytics Hub)
Policy tags
Metadata curation
Administration
Partner Center
Settings
Preview
Release Notes
Search BigQuery resources
Show starred only
Query
Query
Notebooks
Data canvases
Data preparations
Pipelines
External connections
retail_analytics
raw_purchases
Repository
Preview
Store code, edit files, and track changes using version control through repositories or via remote Git based repositories.
*Query
raw_pu... ses
raw_purchases
Query
Open in
Share
Delete
Schema
Details
Insights
Lineage
Data Profile
Data Quality
Field name
Type
Mode
Key
Collation
Default Value
Policy Tags
CustomerID STRING NULLABLE
Age INTEGER NULLABLE
Gender STRING NULLABLE
ItemPurchased STRING NULLABLE
Category STRING NULLABLE
PurchaseAmountUSD NUMERIC NULLABLE
Location STRING NULLABLE
Size STRING NULLABLE
Color STRING NULLABLE
Season STRING NULLABLE
ReviewRating NUMERIC NULLABLE
Edit schema
Job history
Refresh

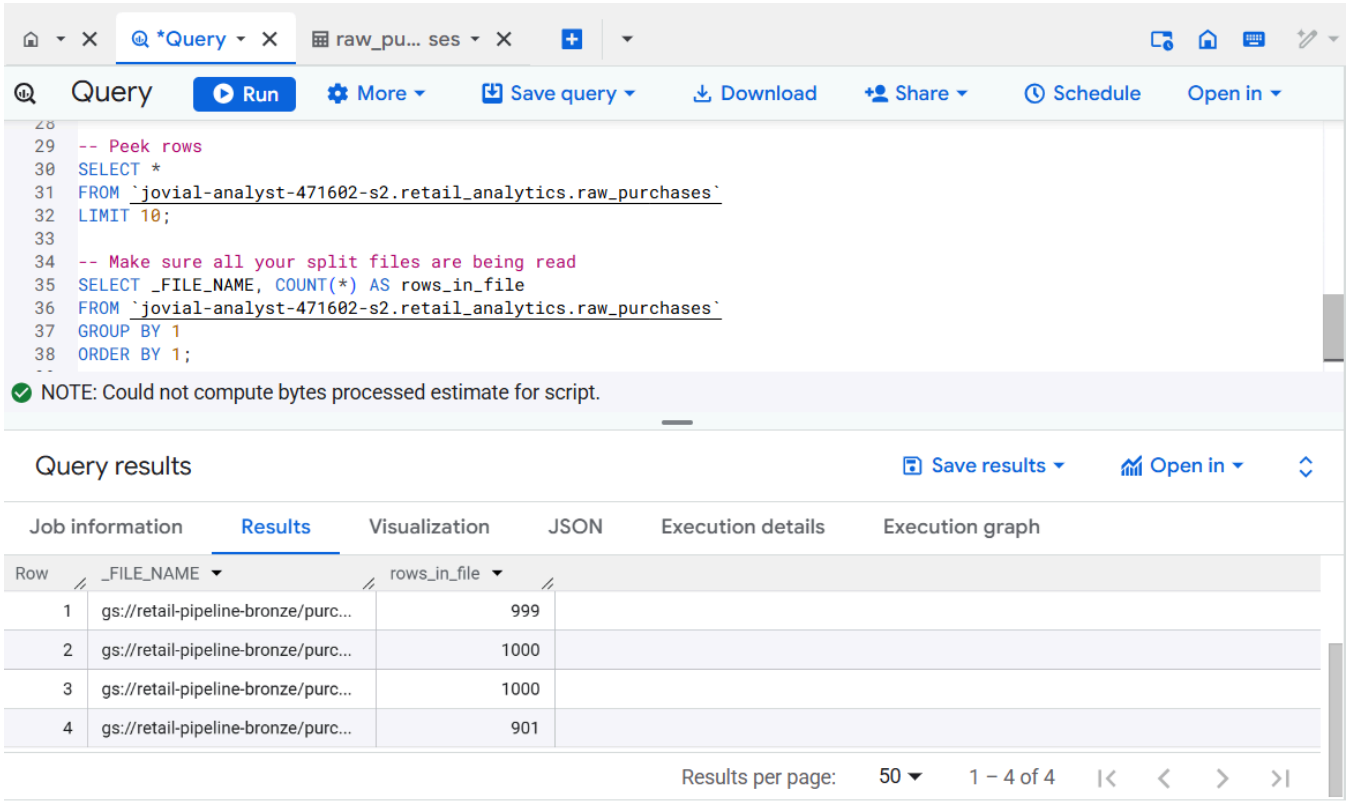
*Query
raw_pu... ses
Query
Run
More
Save query
Download
Share
Schedule
Open in
-- Peek rows
SELECT *
FROM `jovial-analyst-471602-s2.retail_analytics.raw_purchases`
LIMIT 10;
-- Make sure all your split files are being read
SELECT _FILE_NAME, COUNT(*) AS rows_in_file
FROM `jovial-analyst-471602-s2.retail_analytics.raw_purchases`
GROUP BY 1
ORDER BY 1;
NOTE: Could not compute bytes processed estimate for script.
Query results
Save results
Open in
Job information
Results
Visualization
JSON
Execution details
Execution graph
Row
_FILE_NAME
rows_in_file
1 gs://retail-pipeline-bronze/purc... 999
2 gs://retail-pipeline-bronze/purc... 1000
3 gs://retail-pipeline-bronze/purc... 1000
4 gs://retail-pipeline-bronze/purc... 901
Results per page:
50
1 – 4 of 4

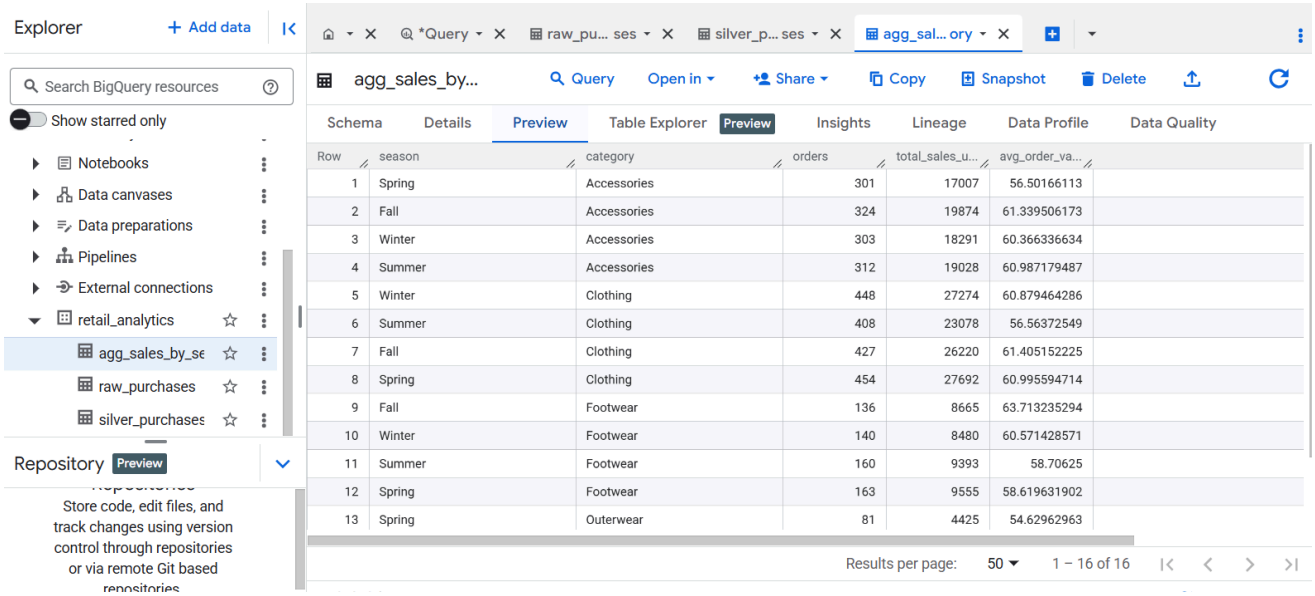
Explorer
+ Add data
Search BigQuery resources
Show starred only
Notebooks
Data canvases
Data preparations
Pipelines
External connections
retail_analytics
agg_sales_by_se
raw_purchases
silver_purchases
Repository
Preview
Store code, edit files, and track changes using version control through repositories or via remote Git based repositories
*Query
raw_pu... ses
silver_p... ses
agg_sal... ory
agg_sales_by...
Query
Open in
Share
Copy
Snapshot
Delete
Schema
Details
Preview
Table Explorer
Preview
Insights
Lineage
Data Profile
Data Quality
Row season category orders total_sales_u... avg_order_va...
1 Spring Accessories 301 17007 56.50166113
2 Fall Accessories 324 19874 61.339506173
3 Winter Accessories 303 18291 60.366336634
4 Summer Accessories 312 19028 60.987179487
5 Winter Clothing 448 27274 60.879464286
6 Summer Clothing 408 23078 56.56372549
7 Fall Clothing 427 26220 61.405152225
8 Spring Clothing 454 27692 60.995594714
9 Fall Footwear 136 8665 63.713235294
10 Winter Footwear 140 8480 60.571428571
11 Summer Footwear 160 9393 58.70625
12 Spring Footwear 163 9555 58.619631902
13 Spring Outerwear 81 4425 54.62962963
Results per page: 50
1 – 16 of 16

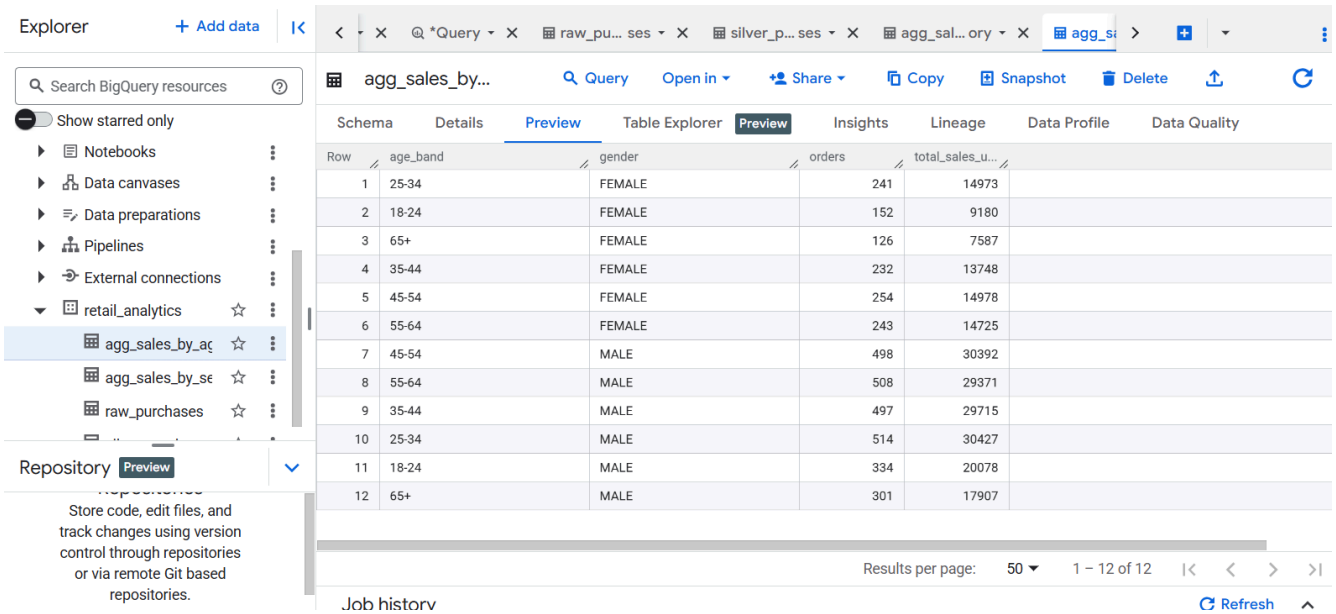
Explorer
+ Add data
Search BigQuery resources
Show starred only
Notebooks
Data canvases
Data preparations
Pipelines
External connections
retail_analytics
agg_sales_by_ag
agg_sales_by_se
raw_purchases
Repository
Preview
Store code, edit files, and track changes using version control through repositories or via remote Git based repositories.
*Query
raw_pu... ses
silver_p... ses
agg_sal... ory
agg_sales_by...
Query
Open in
Share
Copy
Snapshot
Delete
Schema
Details
Preview
Table Explorer
Preview
Insights
Lineage
Data Profile
Data Quality
Row age_band gender orders total_sales_u...
1 25-34 FEMALE 241 14973
2 18-24 FEMALE 152 9180
3 65+ FEMALE 126 7587
4 35-44 FEMALE 232 13748
5 45-54 FEMALE 254 14978
6 55-64 FEMALE 243 14725
7 45-54 MALE 498 30392
8 55-64 MALE 508 29371
9 35-44 MALE 497 29715
10 25-34 MALE 514 30427
11 18-24 MALE 334 20078
12 65+ MALE 301 17907
Results per page: 50
1 – 12 of 12
Job history
Refresh

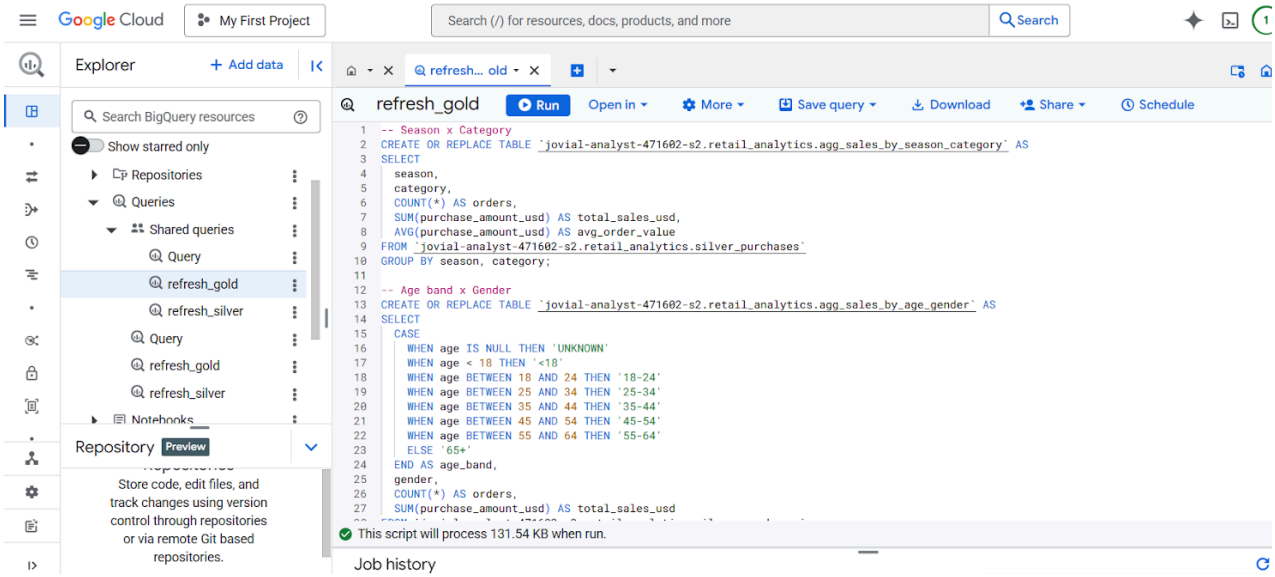

Export as scheduled query Cannot set destination table in jobs with

Details and schedule

Name for scheduled query
refresh_silver

Schedule options

Repeat frequency *
Hours

Repeat frequency *
8 hours

Start now Start at set time

Start date and run time
9/8/25, 11:49 PM EDT

End never Schedule end time

End date EST

Scheduled to run every 8 hours starting now.

Step 5. Visualization & Reporting

Connect the aggregated BigQuery tables to a Looker Studio dashboard. This dashboard provides a visual representation of the processed data, allowing for easy analysis of sales by demographics, season, and product category.